



AWS agents went GA.

The demos are over.

5 gates before production —
and one question at every gate

THE PITCH CHANGED

From "try this model"
to "put agents
to work"

Bedrock grew from
~60 to ~100 models.

Frontier agents run for
hours or days, unattended.

Unattended is where
the bills happen.

THE REAL GA SIGNAL

The real GA signal
isn't a model.

AgentCore harness

idea to running agent, no orchestration

GA · Jun 2026

Evaluations

13 built-in evaluators + Ground Truth

GA · Mar 2026

Policy + Guardrails

fine-grained tool access control

GA · Mar 2026

Markets doing demos don't
ship governance tooling.

GATE 1 + 2

Start smaller than you want to.

One agent.

One tool.

One clear job.

Prove the loop closes on the managed harness before anything fancy.

GATE 3

Evaluate **BEFORE** you scale.

13 built-in evaluators
+ Ground Truth.

Score against expected
answers, not vibes.

If you can't measure it,
you can't run it unattended.

GATE 4

Lock policy while the surface is small.

Tool access on a two-tool agent:
a five-minute job.

Retrofitting it later:
a project.

GATE 5

Deploy against real numbers.

200 TPS

per agent

5,000

sessions · us-east-1 / us-west-2

2,500

sessions · everywhere else

Adjustable defaults — if you
knew to ask before launch.

THE COST RULE

Find the meters before the bill.

model tokens

session time

tool calls

egress

**Two days and ~\$363
taught me this one.**



What did your first agent bill look like?

One number beats any
benchmark I can quote.

Full breakdown →

Featured section on my profile.

Build-to-Learn with AWS GenAI